

CS217 - Data Structures & Algorithm Analysis (DSAA)

Lecture #6

Randomisation & Lower Bounds

Prof. Pietro S. Oliveto

Department of Computer Science and Engineering

Southern University of Science and Technology (SUSTech)

`olivetop@sustech.edu.cn`

<https://faculty.sustech.edu.cn/olivetop>

Reading: Chapters 7.4 and 8.1

➤ Aims of this lecture

- To show how **randomness** can be used in the design of efficient algorithms.
- Glimpse into the **analysis of randomised algorithms**.
- To discuss the class of **comparison sorts**: sorting algorithms that sort by comparing elements.
- To show a general **lower bound** for the running time of a class of sorting algorithms.

➤ A Randomised Version of QuickSort

- Choosing the right pivot element can be tricky – we have no idea *a priori* which pivot elements are good.
- **Solution: leave it to chance!**

RANDOMISED-PARTITION(A, p, r)

1: $i = \text{RANDOM}(p, r)$
2: exchange $A[r]$ with $A[i]$
3: **return** PARTITION(A, p, r)

“Random” picks pivot uniformly at random among all elements.

RANDOMISED-QUICKSORT(A, p, r)

1: **if** $p < r$ **then**
2: $q = \text{RANDOMISED-PARTITION}(A, p, r)$
3: RANDOMISED-QUICKSORT($A, p, q - 1$)
4: RANDOMISED-QUICKSORT($A, q + 1, r$)

➤ Performance of Randomised-QuickSort

- Assume in the following that all elements are distinct.
- What is a worst-case input for Randomised QuickSort?
- **Answer: there is no worst case for Randomised QuickSort!**
- Reason: all inputs lead to the **same runtime behaviour**.
 - The i -th smallest element is chosen with uniform probability.
 - Every split is equally likely, regardless of the input.
 - The runtime is random, but the **random process (probability distribution) is the same** for every input.
- Randomness levels the playing field for all inputs.
 - No one can provide a worst-case input for Randomised-QS.

➤ Runtime of Randomised Algorithms

- For randomised algorithms (in contrast to **deterministic algorithms**) we consider the **expected running time** $E(T(n))$.

- Expectation** of a random variable X :

$$E(X) = \sum_x x \cdot \Pr(X = x)$$

- Example:** for X = roll of fair 6-sided die,

$$E(X) = \sum_x x \cdot \Pr(X = x) = \sum_{x=1}^6 x \cdot \frac{1}{6} = \frac{21}{6} = 3.5$$

- Example** ($X \in \{0, 1\}$): expected #times a coin toss shows heads,

$$E(X) = \sum_x x \cdot \Pr(X = x) = 0 \cdot \Pr(\text{tails}) + 1 \cdot \Pr(\text{heads}) = \Pr(\text{heads}).$$

➤ Linearity of Expectation

- Linearity of expectation:

$$E(X_1 + X_2) = E(X_1) + E(X_2)$$

- Expected number of times 100 coin tosses come up heads:

$$E(X_1 + \dots + X_{100}) = E(X_1) + \dots + E(X_{100}) = 100 \cdot \Pr(\text{heads})$$

- Note: for 0/1-variables the expectation boils down to probabilities.

➤ Number of Comparisons vs. Runtime (1)

For analysing sorting algorithms the **number of comparisons** of elements made is an interesting quantity:

- For QuickSort and other algorithms it can be used as a proxy or substitute for the overall running time (see next slide).
 - Analysing the number of comparisons might be easier than analysing the number of elementary operations.
- **Comparisons can be costly** if the keys to be compared are not numbers, but more complex objects (Strings, Arrays, etc.)
- Algorithms making fewer comparisons might be preferable, even if the overall runtime is the same.
- There is a **lower bound** for the running time of all sorting algorithms that rely on comparisons only.

➤ Number of Comparisons vs. Runtime (2)

- Let $X = X(n)$ be the **number of comparisons** of elements made by QuickSort.
- Comparisons are elementary operations, hence $X(n) \leq T(n)$.
- For each comparison QuickSort only makes $O(1)$ other operations in the for loop.
- Other operations sum to $O(1)$.
- So $X(n) \leq T(n) = O(X(n))$
and thus $T(n) = \Theta(X(n))$
- To show: $E[X(n)] = O(n \log n)$

PARTITION(A, p, r)

```
1:  $x = A[r]$ 
2:  $i = p - 1$ 
3: for  $j = p$  to  $r - 1$  do
4:   if  $A[j] \leq x$  then
5:      $i = i + 1$ 
6:     exchange  $A[i]$  with  $A[j]$ 
7: exchange  $A[i + 1]$  with  $A[r]$ 
8: return  $i + 1$ 
```

Conclusion: we can analyse the **number of comparisons** as a substitute for the runtime in the RAM model.

➤ Expected Time for Randomised-QuickSort

- **Theorem:** the **expected number of comparisons** of Randomised-QuickSort is $O(n \log n)$ for every input where all elements are distinct.
- Proof outline:
 1. Show that here the expectation boils down to probabilities of comparing elements.
 2. Work out the probability of comparing elements.
 3. Putting 1. and 2. together + some maths.
- Follows Section 7.4.2 in the book.

➤ 1. Expectation Boils Down to Probabilities

- For ease of analysis, rename array elements to Z_1, Z_2, \dots, Z_n with $Z_1 < Z_2 < \dots < Z_n$ (hence Z_i is the i -th smallest element)
- **Observation:** each pair of elements is compared at most once.
 - Reason: elements are only compared against the pivot, and after Partition ends the pivot is never touched again.
- Let $X_{i,j}$ be the number of times Z_i and Z_j are compared:

$$X_{i,j} := \begin{cases} 1 & \text{if } z_i \text{ is compared to } z_j \\ 0 & \text{otherwise} \end{cases}$$

- Then the total number of comparisons is $X := \sum_{i=1}^{n-1} \sum_{j=i+1}^n X_{i,j}$
- Taking expectations on both sides and using linearity of expectations:

$$E(X) = E\left(\sum_{i=1}^{n-1} \sum_{j=i+1}^n X_{i,j}\right) = \sum_{i=1}^{n-1} \sum_{j=i+1}^n E(X_{i,j}) = \sum_{i=1}^{n-1} \sum_{j=i+1}^n \Pr(z_i \text{ is compared to } z_j)$$

➤ 2. Probability of comparing Z_i and Z_j

- When is z_i (i -th smallest) compared against z_j (j -th smallest)?
 - If pivot is $x < z_i$ or $z_j < x$ then the decision whether to compare z_i, z_j is **postponed** to a recursive call.
 - If pivot is $x = z_i$ or $x = z_j$ then z_i, z_j **are compared**.
 - If pivot is $z_i < x < z_j$ then z_i and z_j become separated and are **never compared**!
- A decision is only made if $z_i \leq x \leq z_j$. So z_i and z_j are only compared if the **first** pivot chosen amongst $z_i \leq x \leq z_j$ is either z_i or z_j !!
- These are $j - i + 1$ values, out of which 2 lead to z_i, z_j being compared.
- As the pivot element is chosen uniformly at random,

$$\Pr(z_i \text{ is compared to } z_j) = \frac{2}{j - i + 1}$$

- Note: similar numbers are more likely to be compared than dissimilar ones.

➤ 3. Putting things together

$$E(X) = \sum_{i=1}^{n-1} \sum_{j=i+1}^n \Pr(z_i \text{ is compared to } z_j) = \sum_{i=1}^{n-1} \sum_{j=i+1}^n \frac{2}{j-i+1}$$

- Substituting $k := j - i$ yields

$$E(X) = \sum_{i=1}^{n-1} \sum_{k=1}^{n-i} \frac{2}{k+1} \leq 2 \sum_{i=1}^{n-1} \sum_{k=1}^{n-i} \frac{1}{k} \leq 2 \sum_{i=1}^n \sum_{k=1}^n \frac{1}{k} = 2n \sum_{k=1}^n \frac{1}{k}$$

- The sum $\sum_{k=1}^n \frac{1}{k} = \frac{1}{1} + \frac{1}{2} + \frac{1}{3} + \cdots + \frac{1}{n}$

is called **harmonic sum** and is bounded by

$$\sum_{k=1}^n \frac{1}{k} \leq (\ln n) + 1$$

- So we get $E(X) \leq 2n \sum_{k=1}^n \frac{1}{k} = O(n \log n)$

➤ Random Input vs. Randomised Algorithm

- QuickSort is efficient if
 1. The input is random or
 2. The pivot element is chosen randomly
- We have no control over 1., but we can make 2. happen.
- **(Deterministic) QuickSort**
 - **Pro:** the runtime is deterministic for each input
 - **Con:** may be inefficient on some inputs
- **Randomised QuickSort**
 - **Pro:** same behaviour on all inputs
 - **Con:** runtime is random, running it twice gives different times

➤ Other Applications of Randomisation

- **Random sampling**

- Great for big data
- Sample likely reflects properties of the set it is taken from

- **Symmetry breaking**

- Vital for many distributed algorithms

- **Randomised search heuristics**

- General-purpose optimisers, great for complex problems
 - Evolutionary Algorithms / Genetic Algorithms
 - Simulated Annealing
 - Swarm Intelligence
 - Artificial Immune Systems

➤ Summary

- QuickSort has a bad worst-case runtime of $\Theta(n^2)$, but is fast on average.
 - Average-case performance on **random inputs** is $O(n \log n)$.
 - **Randomised QuickSort** sorts any input in **expected time** $O(n \log n)$.
 - Constants hidden in the asymptotic terms are small.
- QuickSort is used in modern programming languages
- **Randomness** can eliminate worst-case scenarios:
 - For randomised QuickSort all inputs are treated the same.
 - The running time is random and can be quantified by considering the **expected running time**: $O(n \log n)$.

► Comparison Sorts

- InsertionSort
- SelectionSort
- MergeSort
- HeapSort
- QuickSort
- All these proceed by comparing elements – we call these **comparison sorts**.
- Sometimes comparisons are the only information available:
 - Multi-dimensional data with no total ordering
(e. g. sorting cars according to speed and price)

► Performance of Comparison Sorts

- The best comparison sorts we have seen so far take time $\Omega(n \log n)$ in the worst case.
- Can we do better?
- Or can we prove that **it's impossible to do better**?
 - Would give us piece of mind (and our boss/customer, ...)
 - Prevents us from wasting time.

► Complexity Theory

(very briefly, more in CS-338 Theory of Computation)

- Complexity theory deals with the **difficulty of problems**.
- **Limits to the efficiency of algorithms**
 - Results like: *every algorithm needs at least time X in the worst case to solve problem Y .*
 - Stops us from **wasting time trying to achieve the impossible!**
 - Informs the design of efficient algorithms.
- Two sides of a coin:

Complexity theory \leftrightarrow Efficient algorithms

► Appetiser: NP-Completeness in a Nutshell

(not relevant for the assessment, but relevant for Computer Science)

- Entscheidungsproblem (decision problem), answer yes/no?
 - **Example:** does there exist an assignment of variables that satisfies a Boolean formula? E.g. $(x_1 \vee \overline{x_2} \vee x_3) \wedge (\overline{x_1} \vee x_4 \vee \overline{x_5}) \wedge \dots$
- NP-complete problems (intuitively, more formal in CS-338)
 - >3000 important problems in different shapes: satisfiability, scheduling, selecting, cutting, routing, packing, colouring, ...
 - It is **easy to verify** that a given solution means “yes”.
 - No one knows how to **find** a solution in polynomial worst-case time!
 - **Either no** NP-complete problem is solvable in polynomial time, **or all of them** are. No one knows! → **“P versus NP problem”**
 - **\$1,000,000 reward** for an answer (let me know if you crack it :-).

► How (Not) to Show Lower Bounds


- How can we show that time $\Theta(\dots)$ is best possible?
- *“We didn’t manage to find a better algorithm.”*
- *“No one in the world has found a better algorithm.”*
 - What if tomorrow someone does?
 - We have to find arguments that apply to **all algorithms that can ever be invented**.
- *“Surely, every efficient algorithm must do things this way.”*
 - You’d be surprised. Efficient algorithms for multiplying matrices start by subtracting elements!

► Comparison Sorts as Decision Trees

- There is one thing that all comparison sorts have to do: **compare elements**!
- Let's strip away all the overhead, data movement, looping, recursing, etc. and take the number of comparisons as lower time bound.
- W.l.o.g. we assume that elements a_1, \dots, a_n are **distinct** – then we can assume that all comparisons have the form $a_i < a_j$.
- A **decision tree** reflects all comparisons **a particular comparison sort** makes, and how the outcome of one comparison determines future comparisons.
 - Like a skeleton of a sorting algorithm.

► Decision tree for a comparison sort

- Inner node $i:j$ means comparing a_i and a_j .




- Leaves: ordering $\pi_1, \pi_2, \dots, \pi_n$ established by the algorithm:


$$a_{\pi_1} \leq a_{\pi_2} \leq \dots \leq a_{\pi_n}$$

A leaf contains a sorted output for a particular input.

- **The execution of a sorting algorithm corresponds to tracing a simple path from the root down to a leaf.**



► Example of a decision tree



► Lower bound for comparison sorts

Theorem: Every comparison sort requires $\Omega(n \log n)$ comparisons in the worst case.

- This includes all comparison sorts that will ever be invented!
- Proof follows; see Theorem 8.1 in the book.
- The theorem can be extended towards an $\Omega(n \log n)$ bound for the **average-case time** (not done here).
- The theorem implies that HeapSort and MergeSort have worst-case time $\Omega(n \log n)$. They are asymptotically **optimal comparison sorts**.

► Proof of the lower bound (1)

- The **worst-case number of comparisons** equals the **length of the longest simple path** from the root to any reachable leaf: we call this the **height h** of the tree (as in HeapSort).
- Every correct algorithm must be able to produce a sorted output for each of the $n!$ possible orderings of the input.
 - \Rightarrow the leaves of the decision tree must be at least $n!$
- A binary tree of height h has no more than 2^h leaves.
 - We'll prove this formally in a bit; let's take this for granted for now.
- To accommodate $n!$ leaves we need $2^h \geq n!$.
- Taking logarithms, this is equivalent to $h \geq \log(n!)$.
- So the worst-case number of comparisons is at least $\log(n!)$.

► What is $\log(n!)$? Proof (2)

- Using $n! \geq \left(\frac{n}{e}\right)^n$ (for $e = \exp(1) = 2.71\dots$) we get

$$\begin{aligned}\log(n!) &\geq \log\left(\left(\frac{n}{e}\right)^n\right) \\ &= n \log(n/e) && (\log(x^y) = y \log(x)) \\ &= n(\log(n) - \log(e)) && (\log(x/y) = \log(x) - \log(y)) \\ &\geq n \log(n) - 1.4427n \\ &= \Omega(n \log n)\end{aligned}$$

- The worst-case number of comparisons is $\Omega(n \log n)$.
- NB for the curious: an average-case bound follows in similar ways as most leaves have to hang at depths of $\Omega(n \log n)$.

► Summary

- **Complexity Theory** gives limits to the efficiency of algorithms.
 - How (not) to prove lower bounds for all algorithms.
- All comparison sorts need time $\Omega(n \log n)$ in the worst case.
 - Decision trees capture the behaviour of every comparison sort.