CS217 - Data Structures & Algorithm Analysis (DSAA)

Lecture #10

▶ AVL Trees: a class of self-balancing trees

Prof. Pietro S. Oliveto

Department of Computer Science and Engineering

Southern University of Science and Technology (SUSTech)

olivetop@sustech.edu.cn

https://faculty.sustech.edu.cn/olivetop

Reading: Lecture notes


▶ Aims of this lecture

- To see a class of **self-balancing trees** guaranteeing operations in time $O(\log n)$.

- To show that the depth of AVL trees is $O(\log n)$.

- To show how to perform insertions and deletions, **rebalancing the tree** through **rotations** whenever it becomes unbalanced.

▶ Self-balancing trees

- There are various types of binary search trees that are guaranteed to have depth $O(\log n)$.
  - AVL Trees
  - 2-3 Trees
  - B-Trees
  - Red-black Trees
  - Splay Trees
  - Van Emde Boas Trees
  - ...

▶ AVL Trees

- Invented by and named after Adelson-Velskii and Landis.

- **Invariant**: all nodes are **locally balanced**.

- A binary tree is called AVL tree if for every node the following holds: the height of the left subtree and the height of the right subtree only differ by at most 1.

- Let v be a node and T_l, T_r be its left and right subtrees, respectively. Then $bal(v) := h(T_l) - h(T_r)$ is the **balance factor** of v, $h(\ )$ denoting the height of a tree.

- In an AVL tree hence for every node v we have $bal(v) \in \{-1, 0, +1\}$.

▶ Balance properties

- The local property does **not** mean that all leaves are on two levels. AVL trees can be lopsided, see this example:

  - However, overall the tree is still pretty balanced.

▶ Minimum number of nodes in an AVL tree

- Let $A(h)$ be the minimum number of nodes in any AVL tree of height h.

  - An AVL tree with height 0 consists of the root only, hence $A(0) = 1$.

  - The smallest AVL tree of height 1 has two nodes, hence $A(1) = 2$.

  - An AVL tree of height h has to have a root with one subtree of height $h - 1$, and the other subtree of height at least $h - 2$.
    Hence $A(h) = 1 + A(h-1) + A(h-2)$.

- This is similar to the Fibonacci numbers (bar the "$1+$"):

  - $Fib(0) = Fib(1) = 1$ and

  - $Fib(h) = Fib(h-1) + Fib(h-2)$.

  - Handy closed form: $$\text{Fib}(k) \geq \frac{1}{\sqrt{5}}\left[\left(\frac{\sqrt{5}+1}{2}\right)^{k+1} - 1\right]$$

▶ Estimating the depth of an AVL tree

Theorem: the height of an AVL tree with n nodes is at most
$$h \leq \frac{1}{\log((\sqrt{5}+1)/2)} \log n \approx 1.44 \log n.$$

- This is only up to 44% deeper than a perfectly balanced tree.

Proof outline:

- Consider the minimum number of nodes in any AVL tree of height h and call it $A(h)$.

  - This means that any AVL tree of height h will have $n \geq A(h)$ nodes.

- Show that $A(h)$ (and thus n) is exponentially large in h.

  - Will show that $A(h)$ is similar to Fibonacci numbers.

- Take logarithms (+maths) to get the claimed bound.

▶ Link to Fibonacci numbers

- We prove by induction that $A(h) = Fib(h+2) - 1$.

- Base case: $A(0) = 1 = 2 - 1 = Fib(2) - 1$
  and $A(1) = 2 = 3 - 1 = Fib(3) - 1$.

- Assume that the claim holds for $A(h-1)$ and $A(h-2)$, then

$$\begin{aligned}
A(h) &= 1 + A(h-1) + A(h-2) & \text{(by recurrence)} \\
&= 1 + \text{Fib}(h+1) - 1 + \text{Fib}(h) - 1 & \text{(2x induction hypothesis)} \\
&= \text{Fib}(h+1) + \text{Fib}(h) - 1 \\
&= \text{Fib}(h+2) - 1 & \text{(by definition of Fib}(h+2)\text{).}
\end{aligned}$$

- Every AVL tree with n nodes and height h has
  $n \geq A(h) \geq Fib(h+2) - 1$.

- Plugging in closed form for Fib gives $\left(\frac{\sqrt{5}+1}{2}\right)^{h+3} \leq \sqrt{5}n + \sqrt{5} + 1$

- Taking logarithm of base $\frac{\sqrt{5}+1}{2}$: $h + 3 \leq \log_{(\sqrt{5}+1)/2}(\sqrt{5}n + \sqrt{5} + 1)$
  $$\Rightarrow h \leq \log_{(\sqrt{5}+1)/2}(n)$$

- Converting to $\log_2$ completes proof.

▶Search in an AVL Tree

- Works like in an ordinary binary search tree.

▶Inserting in an AVL Tree

- Works like in an ordinary binary search tree.

- But the tree may become unbalanced, hence we need to **rebalance**. (We focus on ideas here: code is lab exercise)

- We record the **search path** to new element *z*, and then **work back up the search path** to rebalance **so long as the height of the current subtree has increased**.

- Let *v* be the current node and its **right child *x*** be on the search path (left child is symmetric) -> start at *v* = *z*.parent

▶Insert (1)

Case 1: $bal(v) = 1$.

- Left subtree of *v* was higher than right subtree before insertion.

- After inserting *z*, the right subtree has increased its height, hence the subtree at *v* is now balanced. We set **$bal(v)$ = 0**

- The height of *v* has not changed, hence rebalancing is **done**.

▶Insert (2)

Case 2: $bal(v) = 0$.

- Both subtrees of *v* were balanced before insertion.

- After inserting *z*, the right subtree has increased its height, hence now $bal(v) = -1$.

- The height of the subtree at *v* has **increased** (we cannot stop), hence we need to continue rebalancing at *v*'s parent to check for imbalances further up the tree.

- **If *v* was the root, we stop: done**

▶ Insert (3)

Case 3: $bal(v) = -1$.

- After insertion, the tree has become unbalanced: $bal(v) = -2 \rightarrow$ we need to fix this!

- Search path contains nodes *v, x, w* whose subtrees **increased in height.**

- We distinguish two sub-cases, depending on whether *w* is the right child or the left child of x.

▶ Insert (4)

**Sub-case 3-1**: w is the right child of x.

- The tree is lopsided because of an "outside" problem.

- Now **rotate** the tree to the left: x becomes the parent of *v*, and x's left subtree B becomes a subtree of v. -> **bal(x) = bal(v) = 0**

- Height of whole subtree is the same as before insert. **Done.**

▶ Insert (5)

NB: heights of B and C could be the other way round.

Sub-case 3-2: w is the left child of x.

- The tree is lopsided because of an "inside" problem.

- Now need a **double rotation** to rebalance the tree: a right rotation at x, followed by an immediate left rotation at v.

- **bal(w) = 0**; If after the insertion:

  – bal(w) was -1　　　=> **bal(v) = 1, bal(x) = 0;**

  – bal(w) was 1　　　=> **bal(v)=0,　bal(x) = -1;**

  – bal(w) was 0 *(w = z)* => **bal(v)=0,　bal(x) = 0;**
  　(last case: A,B,C,D all empty)

▶ Insert: Example

▶Insert Rebalancing: Summary

- We go from v=z.parent up the tree until we find a node v with bal(v)=1 coming from right child (-1 coming from left): **Stop** or bal(v)=-1 coming from right (+1 coming from left): **Rotate and Stop**

- If bal(v)=0 **set to** -1 if coming from right child (to +1 if coming from left), and **iterate** unless v is root.

- Rotation:
  - If x=v.right & w=x.right then **L Rotation** (x=v.left & w=x.left => R rot)
  - If x=v.right & w=x.left the **RL Rotation** (x=v.left & w=x.right => LR rot)

▶Runtime of Insert

- Inserting an element takes time $O(h)$.

- Rebalancing:
  - finishes with the first rotation/double rotation.
  - All rotations (L/R/LR/RL) take time $O(1)$.
  - Backing up the search path takes time $O(1)$ for each node on the search path, hence time $O(h)$ overall.
  - This includes the time to update balance factors.

- Total runtime of Insert: $O(h) = O(\log n)$.

▶Deleting in an AVL Tree

- Like for Insert, we work backwards up the search path to rebalance so long as the height of the current subtree has decreased.

- Assume without loss of generality that delete decreased the height of the *left* subtree.

- **Case 1**: $bal(v) = 1$. Here deletion decreased the height of the higher subtree, leading to $\boldsymbol{bal(v) = 0}$.
  However, the height of *v* has decreased, so we need to **iterate** the rebalance procedure with *v*'s parent.

- **Case 2**: $bal(v) = 0$. Then we update $\boldsymbol{bal(v) = -1}$ and note that the height of *v*'s subtree has not decreased, so the rebalancing is complete.

▶Delete (2)

Case 3: $bal(v) = -1$.

- After deletion, the shallower subtree has become even more shallow: $bal(v) = -2$.

- Consider path of nodes *v, x, w* whose subtrees are now too high.

- We distinguish two sub-cases, depending on whether *w* is the right child or the left child of *x*.

▶Delete (3)

Sub-case 3-1: $\text{bal}(x) \in \{-1, 0\}$.

- The tree is lopsided because of an "outside" problem.

- Now **rotate** the tree to the left.

- Two possibilities for the height of B.

- If B was high (bal(x) was 0), => ***bal(v) = -1, bal(x) = 0*** & we're **done**.

- Otherwise, height of the subtree **decreased**, (bal(x) was -1),
  => ***bal(v) = 0, bal(x) = 0*** & **iterate** at x's parent.

▶Delete (4)

Sub-case 3-2: $bal(x) = 1$.

- The tree is lopsided because of an "inside" problem.

- Again need a **double rotation.**

- B and C can have one of two heights; one must be high.

- **Continue** at parent.

▶Runtime of Delete

- Delete may not finish with the first rotation/double rotation.

- Still, the time spent at each node on the search path is $O(1)$, so we still get a time of $O(h) = O(\log n)$.

▶Summary

- AVL trees with n elements have height $O(\log n)$.

- AVL trees with n nodes execute the following operations in time $O(\log n)$

  - **Searching, Minimum, Maximum, Successor**

    - Follows since AVL trees are binary search trees whose height is always $h = O(\log n)$.

  - **Insertion**

  - **Deletion**

- Greater efficiency from a simple idea: rotating nodes.